

Cloudformation Template Documentation

Overview

This document will provide a brief explanation of the deployed environment and a brief explanation of the CloudFormation templates.

Additional material

There is a high-level diagram which provides an overview of the solution in the same location as this document, the name of the file is *overview.pdf*

Deployment Strategy

I used CloudFormation templates to deploy the solution as it provides potential for further automation.

I tried to logically separate each part of the solution as it allowed for faster deployments and easier troubleshooting with the intention being to combine all the yml files into one larger all-encompassing CloudFormation template once everything was working.

The repo consists of the following templates:

vpc-m.yml

This is an aws template that I used to deploy a custom VPC, found here:

<https://docs.aws.amazon.com/codebuild/latest/userguide/cloudformation-vpc-template.html>

Only changes I made were to the CIDR blocks via the parameters, this can easily be stripped out so that no user interaction would be required.

s3.yml

For the creation of an s3 bucket for use by the elastic beanstalk deployment.

I enabled versioning and encryption, not publicly viewable by default.

roles.yml

This was used for troubleshooting while running up the elastic beanstalk stack. I ran into quite a few issues and so I created a new role and allowed it to assume

arn:aws:iam::aws:policy/AWSElasticBeanstalkFullAccess while it spun up resources for the beanstalk deployment. I eventually discovered that the issue was the method I was using to select instance types for the autoscaling group. I added some additional *OptionName* parameters in eb-m.yml which allowed it proceed further into the deployment.

rds-m.yml

For the deployment of the RDS Postgres solution into a private subnet of the custom VPC via providing a DBSubnetGroup, which requires subnets in at least two unique Availability Zones.

Deployed postgres 11.6 and allocated 20GB of storage.

eb-m.yml

For the elastic beanstalk deployment, unfortunately beanstalk likes to deploy into the default VPC, therefore I needed to add the following *OptionNames*:

- VPCID
- Subnets
- ELBSubnets -
- DBSubnets
- ELBScheme

The elastic beanstalk deployment also pulls the ruby app from s3.

Testing

Was performed manually, will briefly outline some of the tests below:

Database

Spun up an ec2 instance which I used to connect to the postgres db

```
[ec2-user@ip-10-10-10-52 ~]$ psql -U sysalykes -h mablerdsm.cbdlh26f7kda.ap-southeast-2.rds.amazonaws.com -d mableDBm
```

Password for user sysalykes:

psql (9.2.24, server 11.6)

WARNING: psql version 9.2, server version 11.0.

Some psql features might not work.

SSL connection (cipher: ECDHE-RSA-AES256-GCM-SHA384, bits: 256)

Type "help" for help.

```
mableDBm=> \l
```

List of databases

Name	Owner	Encoding	Collate	Ctype	Access privileges
mableDBm	sysalykes	UTF8	en_US.UTF-8	en_US.UTF-8	
postgres	sysalykes	UTF8	en_US.UTF-8	en_US.UTF-8	

```

rdsadmin | rdsadmin | UTF8 | en_US.UTF-8 | en_US.UTF-8 | rdsadmin=CTc/rdsadmin
template0 | rdsadmin | UTF8 | en_US.UTF-8 | en_US.UTF-8 | =c/rdsadmin +
          |          |          |          |          | rdsadmin=CTc/rdsadmin
template1 | sysalykes | UTF8 | en_US.UTF-8 | en_US.UTF-8 | =c/sysalykes +
          |          |          |          |          | sysalykes=CTc/sysalykes
(5 rows)

mableDBm=> \q

```

Elastic Beanstalk

awseb-e-hamrprckkg-stack

Delete

Update

Stack actions

Stack info

Events

Resources

Outputs

Parameters

Template

Change sets

Outputs (1)

Q Search outputs

Key ▲	Value ▼	Description
AWSEBLoadBalancerURL	http://awseb-e-h-AWSEBLoa-A1AYINKQID3A-94476892.ap-southeast-2.elb.amazonaws.com	The ElasticBeanstalk ELB URL of the website

mabl-node-l2p1btfwoe1u.eba-2p2td52m.ap-southeast-2.elasticbeanstalk.com/index.html

Congratulations

Your first AWS Elastic Beanstalk **Ruby** Application is now running on your own dedicated environment in the AWS Cloud

This environment is launched with Elastic Beanstalk Ruby Platform

What's Next?

- [AWS Elastic Beanstalk overview](#)
- [AWS Elastic Beanstalk concepts](#)
- [Deploy a Ruby on Rails Application to AWS Elastic Beanstalk](#)
- [Deploy a Sinatra Application to AWS Elastic Beanstalk](#)
- [Customizing and Configuring a Ruby Container](#)
- [Working with Logs](#)

Cloudformation Stacks

Stacks (5)

Q Filter by stack name

Active ▼

View nested

Refresh

Delete

Update

	Stack name	Status	Created time ▼	Description
<input type="radio"/>	awseb-e-hamrprckkg-stack	✔ CREATE_COMPLETE	2020-09-05 17:47:10 UTC+1000	AWS Elastic Beanstalk environment (Name: 'mabl-node-L2P1BTFW0E1U' Id: 'e-hamrprckkg')
<input type="radio"/>	mable-ebm	✔ CREATE_COMPLETE	2020-09-05 17:46:55 UTC+1000	-
<input type="radio"/>	mable-rds	✔ CREATE_COMPLETE	2020-09-05 17:04:27 UTC+1000	Description: AWS CloudFormation Sample Template for creating an Amazon RDS DB instance:
<input type="radio"/>	mable-vpc	✔ CREATE_COMPLETE	2020-09-05 16:33:44 UTC+1000	This template deploys a VPC, with a pair of public and private subnets spread across two Availabi
<input type="radio"/>	mable-s3	✔ CREATE_COMPLETE	2020-09-05 10:34:26 UTC+1000	CloudFormation template for s3 bucket

Further CloudFormation Optimisations

- Once the VPC is created, it is then required to manually replace any hardcoded values for VPCIDs and Subnets in the cloudformation templates.
- Create a single cloudformation template for the entire stack.